

# Musings on data access

Colin Bradford  
[cjbradford@gmail.com](mailto:cjbradford@gmail.com)

# Problem

- Data is relational
- SQL lets us describe the relations
- But the results are a rectangular table

# Example

- Journey has multiple stops
- Journey has multiple passengers
- Joining both tables gives stops x passengers rows
- Need to deduplicate the data

# Possibilities

- Multiple queries
- UNION queries with nulls
- Others?

# Alternative – describe the result

```
{
  "name" : "trip",
  "index" : {
    "name" : "trip_id_index",
    "in" : [ 5744 ]
  },
  "join" : [
    {
      "name" : "stops",
      "index" : {
        "name" : "PRIMARY",
        "parent_column" : "trip_id"
      }
    },
    {
      "name" : "passengers",
      "index" : {
        "name" : "PRIMARY",
        "parent_column" : "trip_id"
      }
    }
  ]
}
```

# Which gives

```
[
  {
    trip_id => 5744,
    customer_id => 6455,
    stops => [
      { address => 'BBC, W1A 1AA', },
      { address => 'Euston Station', },
    ],
    passengers => [
      { name => 'Mr J. Smith', },
      { name => 'Mr A. N. Other', }
    ]
  },
]
```

# Or (faster over the wire)

```
{  
  data => [  
    { trip_id => 5744, customer_id => 6455, },  
  ],  
  stops => {  
    5744 => [  
      { address => 'BBC, W1A 1AA', },  
      { address => 'Euston Station', }  
    ]  
  },  
  passengers => {  
    5744 => [  
      { name      => 'Mr J. Smith', },  
      { name      => 'Mr A. N. Other', }  
    ]  
  },  
}
```

# Implementation

- Do joins in Perl
  - which means they run on the client, not the server
- Can target multiple back end data stores
  - SQL
  - Redis
  - MongoDB
  - CouchDB
  - HandlerSocket



# HandlerSocket?

- Access into the InnoDB storage engine of MySQL
- Lightweight protocol
- Can only look up on indexed columns
- Fast
- But Perl interface slows it down

# So, does it work?

- Yes (ish), for reads only
- Initial implementation on GitHub:  
<https://github.com/cjbradford/DataKeyValue>
- Not performance optimised

# Advantages

- Can join across data stores
- Can join with data stores that don't support joins
- Can layer: Add a shard layer over another backend
- Load appears to move from the data store to the client

# Performance

- Measured using Benchmark.pm
- Results open to interpretation
- Data set is as similar as I can make it – not optimised for Key/Value data store
- Data stores running on the same machine as the client
- Single SATA disk, 8Gb ram, i5-2500K CPU @3.3GHz

# Query with a join

```
SELECT p.chanid, p.starttime, p.endtime, p.title, p.subtitle,  
p.description, c.chanid,  
c.channum, c.callsign, c.sourceid, c.serviceid, c.mplexid  
FROM program p  
LEFT JOIN channel c ON p.chanid = c.chanid  
WHERE starttime >= ? AND starttime <= ?  
(123 rows back)
```

CouchDB:	127	wall	(	4.78	usr	0.24	sys)	@	4.98/s	(n=25)
MongoDB:	8	wall	(	5.13	usr	0.68	sys)	@	65.92/s	(n=383)
Redis:	6	wall	(	4.35	usr	0.88	sys)	@	138.62/s	(n=725)
HSocket:	8	wall	(	5.32	usr	0.08	sys)	@	375.19/s	(n=2026)
PerlMyISAM:	5	wall	(	5.27	usr	0.03	sys)	@	543.77/s	(n=2882)
PerlInnoDB:	5	wall	(	5.05	usr	0.00	sys)	@	546.53/s	(n=2760)
SQLMyISAM:	6	wall	(	5.12	usr	0.09	sys)	@	1025.91/s	(n=5345)
SQLInnoDB:	5	wall	(	5.14	usr	0.02	sys)	@	1035.85/s	(n=5345)

# Single table, lots of rows

```
SELECT p.chanid, p.starttime, p.endtime, p.title, p.subtitle,  
p.description  
FROM program p  
WHERE starttime BETWEEN ? AND ?  
(123 rows)
```

CouchDB:	50	wall	(	5.01	usr	0.36	sys)	@	83.24/s	(n=447)
MongoDB:	5	wall	(	4.97	usr	0.03	sys)	@	383.40/s	(n=1917)
Redis:	6	wall	(	5.10	usr	0.18	sys)	@	410.61/s	(n=2168)
PerlMyISAM:	5	wall	(	5.15	usr	0.00	sys)	@	1243.69/s	(n=6405)
PerlInnoDB:	6	wall	(	5.21	usr	0.02	sys)	@	1247.23/s	(n=6523)
SQLInnoDB:	5	wall	(	5.28	usr	0.03	sys)	@	1251.60/s	(n=6646)
SQLMyISAM:	5	wall	(	5.23	usr	0.07	sys)	@	1255.85/s	(n=6656)
HSocket:	12	wall	(	6.27	usr	0.26	sys)	@	1365.08/s	(n=8914)

# Single table, one row

```
SELECT c.chanid, c.channum, c.callsign, c.sourceid, c.serviceid,  
c.mplexid  
FROM channel c  
WHERE chanid BETWEEN ? AND ?  
(1 row)
```

CouchDB:	133	wall	(	4.80	usr	0.25	sys)	@	327.72/s	(n=1655)
MongoDB:	7	wall	(	5.24	usr	0.56	sys)	@	8506.72/s	(n=49339)
Redis:	6	wall	(	4.06	usr	1.18	sys)	@	12241.41/s	(n=64145)
SQLInnoDB:	9	wall	(	6.77	usr	0.72	sys)	@	12926.17/s	(n=96817)
PerlMyISAM:	6	wall	(	4.60	usr	0.53	sys)	@	15192.98/s	(n=77940)
PerlInnoDB:	6	wall	(	4.86	usr	0.37	sys)	@	15373.23/s	(n=80402)
SQLMyISAM:	7	wall	(	5.61	usr	0.56	sys)	@	15691.57/s	(n=96817)
HSocket:	7	wall	(	3.77	usr	1.48	sys)	@	30631.43/s	(n=160815)

# Single table, PK lookup

```
SELECT c.chanid, c.channum, c.callsign, c.sourceid, c.serviceid,  
c.mplexid  
FROM channel c  
WHERE chanid = ?  
(1 row)
```

CouchDB:	136	wall	(	4.89	usr	0.36	sys)	@	644.19/s	(n=3382)
MongoDB:	7	wall	(	4.33	usr	0.92	sys)	@	12161.52/s	(n=63848)
SQLMyISAM:	6	wall	(	5.24	usr	0.49	sys)	@	16896.51/s	(n=96817)
PerlMyISAM:	7	wall	(	4.78	usr	0.47	sys)	@	17568.57/s	(n=92235)
PerlInnoDB:	5	wall	(	4.64	usr	0.36	sys)	@	17757.20/s	(n=88786)
SQLInnoDB:	6	wall	(	4.88	usr	0.36	sys)	@	18476.53/s	(n=96817)
Redis:	6	wall	(	3.98	usr	1.19	sys)	@	23353.97/s	(n=120740)
HSocket:	7	wall	(	3.55	usr	1.57	sys)	@	29809.77/s	(n=152626)



# Closing thoughts

- So far, idea isn't obviously broken
- In most of the applications I work with:
  - Most writes are to a single “table”
  - Transactions can be worked around
  - Moving load to clients is worthwhile

Thank you